

# assignment 12

June 3, 2023

## 0.1 Assignment 12

```
[ ]: import numpy as np
from keras.datasets import mnist
from keras.models import Model
from keras.layers import Conv2D, Flatten, Dense, Input, Lambda, Reshape,
    ↪Conv2DTranspose, Layer
from keras import metrics
from keras import backend as K
K.clear_session()
import tensorflow._api.v2.compat.v1 as tf
tf.disable_v2_behavior()
import matplotlib.pyplot as plt
```

```
[ ]: img_shape = (28, 28, 1)
batch_size = 16
latent_dim = 2

input_img = Input(shape=img_shape)

x = Conv2D(32, 3, padding='same', activation='relu')(input_img)
x = Conv2D(64, 3, padding='same', activation='relu', strides=(2, 2))(x)
x = Conv2D(64, 3, padding='same', activation='relu')(x)
x = Conv2D(64, 3, padding='same', activation='relu')(x)

shape_before_flattening = K.int_shape(x)

x = Flatten()(x)
x = Dense(32, activation='relu')(x)

z_mean = Dense(latent_dim)(x)
z_log_var = Dense(latent_dim)(x)
```

```
[ ]: def sampling(args):
    z_mean, z_log_var = args
    epsilon = K.random_normal(shape=(K.shape(z_mean)[0], latent_dim), mean=0.,
    ↪stddev=1.)
    return z_mean + K.exp(z_log_var) * epsilon
```

```
z = Lambda(sampling)([z_mean, z_log_var])
```

```
[ ]: decoder_input = Input(K.int_shape(z)[1:])

x = Dense(np.prod(shape_before_flattening[1:]),
          activation='relu')(decoder_input)
x = Reshape(shape_before_flattening[1:])(x)
x = Conv2DTranspose(32, 3, padding='same', activation='relu', strides=(2, 2))(x)
x = Conv2D(1, 3, padding='same', activation='sigmoid')(x)

decoder = Model(decoder_input, x)

z_decoded = decoder(z)
```

```
[ ]: class CustomVariationalLayer(Layer):

    def vae_loss(self, x, z_decoded):
        x = K.flatten(x)
        z_decoded = K.flatten(z_decoded)
        xent_loss = metrics.binary_crossentropy(x, z_decoded)
        kl_loss = -5e-4 * K.mean(
            1 + z_log_var - K.square(z_mean) - K.exp(z_log_var), axis=-1)
        return K.mean(xent_loss + kl_loss)

    def call(self, inputs):
        x = inputs[0]
        z_decoded = inputs[1]
        loss = self.vae_loss(x, z_decoded)
        self.add_loss(loss, inputs=inputs)
        return x

y = CustomVariationalLayer()([input_img, z_decoded])
```

```
[ ]: vae = Model(input_img, y)
vae.compile(optimizer='rmsprop', loss=None)
vae.summary()

(x_train, _), (x_test, y_test) = mnist.load_data()

x_train = x_train.astype('float32') / 255.
x_train = x_train.reshape(x_train.shape + (1,))
x_test = x_test.astype('float32') / 255.
x_test = x_test.reshape(x_test.shape + (1,))
```

```
vae.fit(x=x_train, y=None, shuffle=True, epochs=10, batch_size=batch_size,
      validation_data=(x_test, None))
```

WARNING:tensorflow:Output custom\_variational\_layer missing from loss dictionary.  
We assume this was done on purpose. The fit and evaluate APIs will not be  
expecting any data to be passed to custom\_variational\_layer.

Model: "model\_1"

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	[(None, 28, 28, 1)]	0	[]
conv2d (Conv2D) ['input_1[0][0]']	(None, 28, 28, 32)	320	
conv2d_1 (Conv2D) ['conv2d[0][0]']	(None, 14, 14, 64)	18496	
conv2d_2 (Conv2D) ['conv2d_1[0][0]']	(None, 14, 14, 64)	36928	
conv2d_3 (Conv2D) ['conv2d_2[0][0]']	(None, 14, 14, 64)	36928	
flatten (Flatten) ['conv2d_3[0][0]']	(None, 12544)	0	
dense (Dense) ['flatten[0][0]']	(None, 32)	401440	
dense_1 (Dense)	(None, 2)	66	['dense[0][0]']
dense_2 (Dense)	(None, 2)	66	['dense[0][0]']
lambda (Lambda) ['dense_1[0][0]', 'dense_2[0][0]']	(None, 2)	0	
model (Functional) ['lambda[0][0]']	(None, 28, 28, 1)	56385	
custom_variational_layer (CustomVariationalLayer)	(None, 28, 28, 1)	0	'model[0][0]'

=====

Total params: 550,629

Trainable params: 550,629

Non-trainable params: 0

-----  
-----

Train on 60000 samples, validate on 10000 samples

Epoch 1/10

59952/60000 [=====>.] - ETA: 0s - loss: 342.0744

c:\ProgramData\Anaconda3\lib\site-packages\keras\engine\training\_v1.py:2045:

UserWarning: `Model.state\_updates` will be removed in a future version. This property should not be used in TensorFlow 2.0, as `updates` are applied automatically.

updates = self.state\_updates

60000/60000 [=====] - 64s 1ms/sample - loss: 341.8009 -  
val\_loss: 0.1992

Epoch 2/10

60000/60000 [=====] - 57s 943us/sample - loss: 0.1945 -  
val\_loss: 0.1901

Epoch 3/10

60000/60000 [=====] - 89s 1ms/sample - loss: 0.1897 -  
val\_loss: 0.1879

Epoch 4/10

60000/60000 [=====] - 58s 962us/sample - loss: 0.1871 -  
val\_loss: 0.1877

Epoch 5/10

60000/60000 [=====] - 95s 2ms/sample - loss: 0.1853 -  
val\_loss: 0.1836

Epoch 6/10

60000/60000 [=====] - 54s 898us/sample - loss: 0.1840 -  
val\_loss: 0.1842

Epoch 7/10

60000/60000 [=====] - 64s 1ms/sample - loss: 0.1829 -  
val\_loss: 0.1833

Epoch 8/10

60000/60000 [=====] - 93s 2ms/sample - loss: 0.1821 -  
val\_loss: 0.1821

Epoch 9/10

60000/60000 [=====] - 61s 1ms/sample - loss: 0.1815 -  
val\_loss: 0.1820

Epoch 10/10

60000/60000 [=====] - 67s 1ms/sample - loss: 0.1809 -  
val\_loss: 0.1817

[ ]: <keras.callbacks.History at 0x2555e3ccd30>

```

[ ]: import matplotlib.pyplot as plt
from scipy.stats import norm
import os
filepath = 'results/vae/'
file = "MNIST_VAE_grid.png"
if not os.path.exists(filepath):
    os.makedirs(filepath)

n = 15
digit_size = 28
figure = np.zeros((digit_size * n, digit_size * n))
grid_x = norm.ppf(np.linspace(0.05, 0.95, n))
grid_y = norm.ppf(np.linspace(0.05, 0.95, n))

for i, yi in enumerate(grid_x):
    for j, xi in enumerate(grid_y):
        z_sample = np.array([[xi, yi]])
        z_sample = np.tile(z_sample, batch_size).reshape(batch_size, 2)
        x_decoded = decoder.predict(z_sample, batch_size=batch_size)
        digit = x_decoded[0].reshape(digit_size, digit_size)
        figure[i * digit_size: (i + 1) * digit_size,
              j * digit_size: (j + 1) * digit_size] = digit

plt.figure(figsize=(10, 10))
plt.imshow(figure, cmap='Greys_r')
plt.savefig(os.path.join(filepath, file))
plt.show()

```

